

CS222/CS122C: Principles of Data Management

UCI, Fall 2019
Notes #04

Schema versioning and
File organizations

Instructor: Chen Li

Schema Versioning

- ❖ How to handle existing records when the schema is changed?
- ❖ Scheme versioning technique adds the version of the schema to each record
 - ❖ All versions of the schema are kept in the catalog
 - ❖ When the schema changes, create a new schema version
 - ❖ Records are interpreted based on it's schema version and current schema during a query

Table	Record	Version
	a1, b1, c1	1
	a2, b2, c2	1
	a3, b3, c3, d3	2
	a4, b4, d4	3

Version	Schema
1	A, B, C
2	A, B, C, D
3	A, B, D

Catalog

Example

- ❖ Create table R(A, B, C)
- ❖ Insert 1 million records
 - All records are with schema version 1
- ❖ Add attribute D
 - Create new schema version 2
- ❖ Insert 10 records
 - 10 records are with schema version 2
- ❖ Select * from R
 - records with version 1 are padded with null for field D
- ❖ Drop attribute D
 - Create new schema version 3
- ❖ Select * from R
 - field D are truncated from records with version 2

Next topic: File Organizations

Many alternatives exist. *Each one is ideal for some situations, but not so good in others:*

- Heap (random ordered) files: Suitable when typical access is a file scan retrieving all record or access comes through a variety of secondary indexes.
- Sorted Files: Best if records must be retrieved in some order, or only a 'range' of records is needed.
- Indexes: Data structures to organize records via trees or hashing.
 - Like sorted files, they speed up searches for a subset of records, based on values in certain ("search key") fields.
 - Updates are much faster than in sorted files.

Cost Model

We will ignore CPU costs, for simplicity, so:

- **B:** The number of data pages
- **R:** Number of records per page
- **D:** (Average) time to read or write disk page
- Counting the number of page I/Os ignores gains of prefetching a sequence of pages; thus, even the real I/O cost is only roughly approximated for now.
- Average-case analysis; based on several simplistic assumptions.

** Good enough to convey the overall trends!*

Comparison of File Organizations

- ❖ Heap files (random order; insert at eof)
- ❖ Sorted files, sorted on $\langle \text{age}, \text{sal} \rangle$

Operations to Compare

- ❖ Scan: Fetch all records from disk
- ❖ Equality search
- ❖ Range selection
- ❖ Insert a record
- ❖ Delete a record

Assumptions for Our Analysis

❖ Heap Files:

- Equality selection on key; exactly one match.

❖ Sorted Files:

- File compacted after a deletion (*vs.* a deleted bit).

Cost of Operations

	(a) Scan	(b) Equality	(c) Range	(d) Insert	(e) Delete
(1) Heap					
(2) Sorted					

** Several assumptions underlie these (rough) estimates!*

Scan: scan all the records in the file

Equality: like "id = 123", suppose only one record satisfy the equality

Range: like "5 < id < 100"

Insert:

Delete:

Assume # of pages in the file: B

Cost of Operations (disk IOs)

assume no cost to locate page with enough free space

	(a) Scan	(b) Equality / Search	(c) Range	(d) Insert	(e) Delete
(1) Heap	BD	0.5BD record distributed uniformly	BD <i>locate id = t</i>	2D read page and flush page back	Search + D
(2) Sorted T(id, name, sal,...) Sorted only based on one attribute	BD	D log 2 B	D(log 2 B + # pgs with match recs) <i>scan from id = 5 to id = 100</i>	Search + BD <u>B/2*d + B/2*D</u> <i>shift all the following page</i>	Search + BD

Get benefit only for the sorted attribute

* Several assumptions underlie these (rough) estimates!

B: total num of records
D: cost of operation for one disk IO

Delete of heap file:

- 1) given a RID -> 2D(read and write)
- 2) given a value -> search + D(write back)

Average shift B/2 pages
-> read to memory then write back to disk