

# CS222/CS122C: Principles of Data Management

UCI, Fall 2019  
Notes #05

## Index Overview and ISAM Tree Index

Instructor: Chen Li

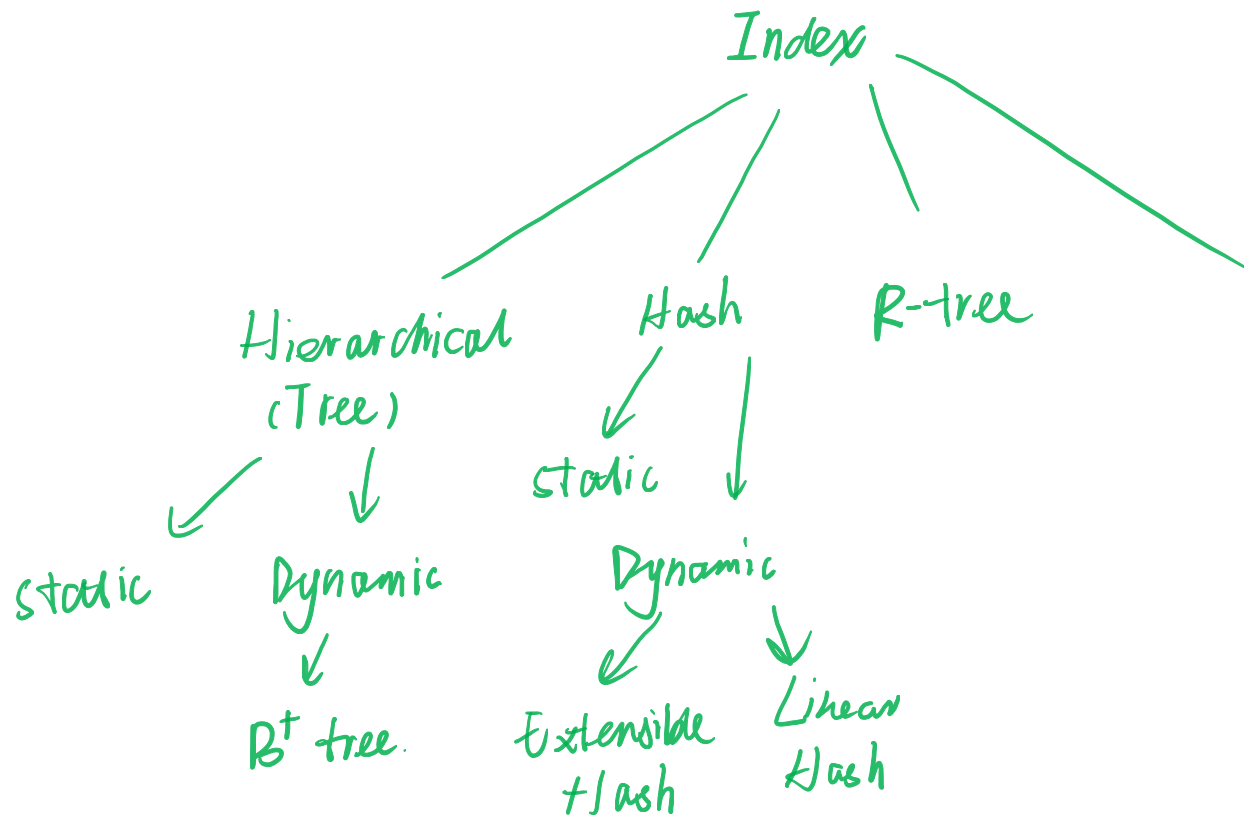
# *Index families*

## ❖ Tree-based indexes

- ISAM: static structure (Indexed Sequential Access Method)
- B+ tree: dynamic, adjusts gracefully under inserts and deletes.

## ❖ Hash tables

- Static hashing
- Dynamic hashing



# Indexes


- ❖ An index on a file speeds up selections on the *search key fields* for the index.
  - Any subset of the fields of a relation can serve as the search key for an index on the relation.
  - Search key is **not** the same as a *key* (a minimal set of fields that uniquely identify a record in a relation).

Allow same  
Search key  
Allow duplicate


# *Data Entries*

- ❖ An index contains a collection of *data entries*, and supports efficient retrieval of all data entries  $k^*$  with a given key value  $k$ .
  - Given data entry  $k^*$ , we can find one record with key  $k$  with just more disk I/O.

*E.g., Tree index*



- ❖ Leaf pages contain data entries, and are chained (prev & next)
- ❖ Non-leaf pages have index entries; used only to direct searches:



# Alternatives for Data Entry $k^*$ in Index

❖ In a data entry  $k^*$  we can store:

- Actual data record with key value  $k$ , or
- $\langle k, \text{rid of data record with search key value } k \rangle$ , or
- $\langle k, \text{list of rids of data records with search key } k \rangle$

*clustered tree*

❖ Choice of alternative for data entries is orthogonal to the indexing technique used to locate data entries with a given key value  $k$ .

- Examples of indexing techniques: B+ trees, hash-based structures, R trees, ...
- Index's job: direct searches to desired data entries
- Rid alternative in secondary indexes: primary key

*Each Rid is unique  $\rightarrow$  primary key*

# Alternatives for Data Entries (Contd.)

## ❖ Alternative 1: Data Records Live in Index

- If this is used, index structure is actually a file organization for the data records (instead of a Heap file or a sorted or hashed file).
- At most one index on a given collection of data records can use Alternative 1. (Otherwise, data records are duplicated, leading to redundant storage and potential inconsistency. Right?)
- If data records are very large, # of (leaf) pages containing data entries is high. Implies size of auxiliary information in the index is also large, typically.

↓ 否则,  
有几个树  
就有几个  
拷贝



# *Alternatives for Data Entries (Contd.)*

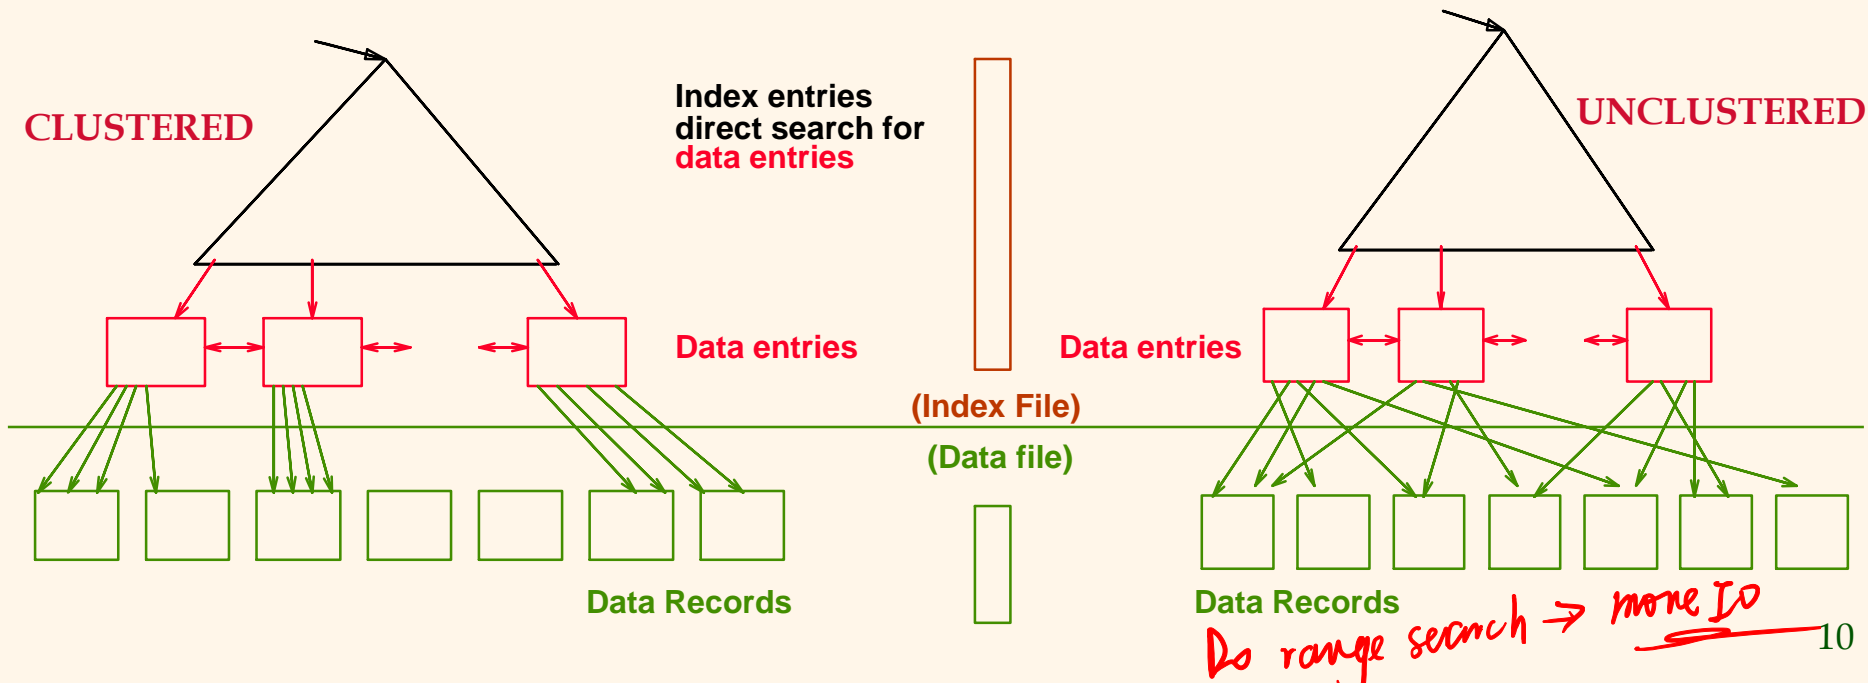
- ❖ Alternatives 2 and 3: Key/Rid or Key/RidList
  - Data entries typically much smaller than data records. (Portion of index structure used to direct searches, which depends on size of data entries, is much smaller than with Alternative 1.)
  - Alternative 3 more compact than Alternative 2, but leads to variable sized data entries even if the search keys are of fixed length.
  - Can treat Key/Rid pair in a composite key-like fashion in higher levels of index to handle case where a (big) RidList could overflow a leaf page.

# Index Classification

- ❖ *Primary vs. secondary*: If search key contains the primary key, then called the primary index.
  - *Unique* index: Search key contains a *candidate* key.
- ❖ *Clustered vs. unclustered*: If order of data records is the same as, or 'close to', the order of stored data entries, then called a clustered index.
  - Alternative 1 implies clustered; in practice, clustered also implies Alternative 1 (sorted files are rare).
  - A file can be clustered on at most one search key.
  - Cost of retrieving data records via index varies *greatly* based on whether index is clustered or not!

# Clustered vs. Unclustered Index

- ❖ Suppose that Alternative (2) is used for data entries, and that the data records are stored in a Heap file.
  - To build a clustered index, first sort the Heap file (with some free space left on each page for future inserts).
  - Overflow pages may be needed for inserts. (Thus, order of data recs is 'close to', but not identical to, the sort order.)

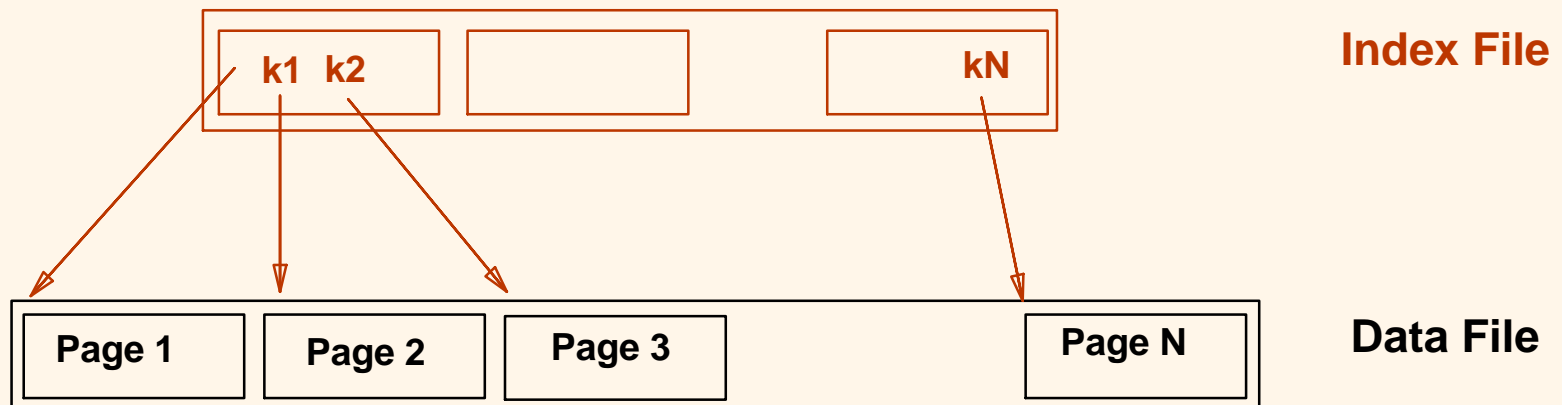


## *Next: Tree Indexes*

- ❖ ISAM: static structure (Indexed Sequential Access Method)
- ❖ B+ tree: dynamic, adjusts gracefully under inserts and deletes.


# Range Searches

- ❖ *‘Find all students with  $gpa > 3.0$ ’*
  - If data is in sorted file, do binary search to find first such student, then scan to find others.
  - Cost of binary search can be quite high.
- ❖ Simple idea: Create an ‘index’ file.

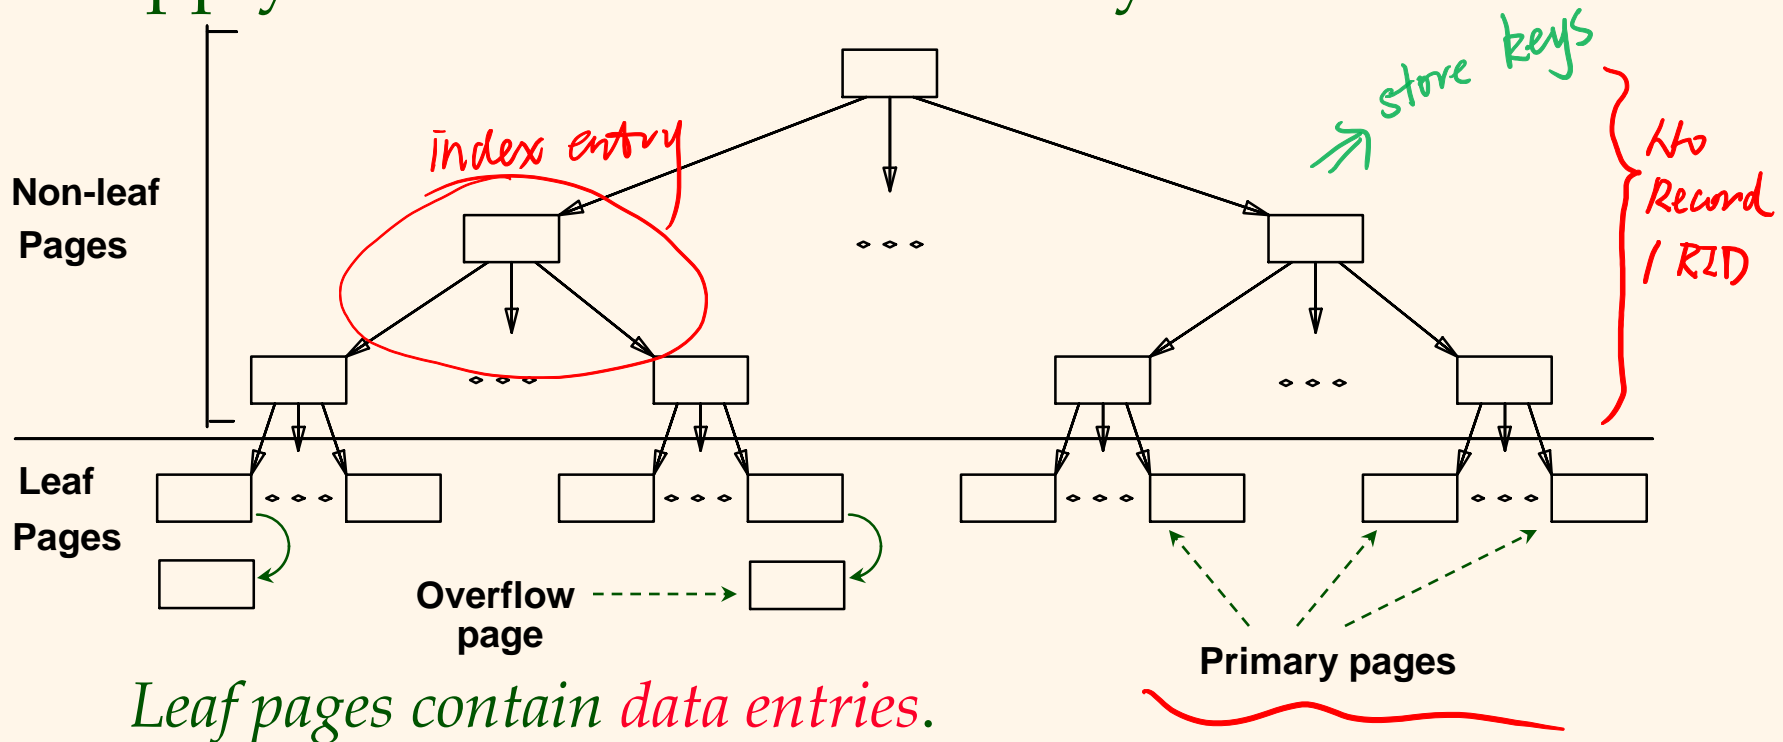


*Can now do the binary search on the (smaller) index file!*

# ISAM



- ❖ Index file may still be quite large. But, we can apply the same idea recursively to address that!



Index key attribute: price

$k^*$  →  
↓  
key entry

1) Record itself = clustered Index

2) RID for  $k$ : much smaller



(k, RID) ↓ heap file is the data home, could have multiple index tree.

3) List of RIDs

↓  
Duplicate

↓  
for the records which have same key

then tree is the data home, the usually only have one copy




← pointer to subtree

$k_1 \leq x < y < k_2$

只有节点可以 equal

# Comments on ISAM

- ❖ *File creation:* Leaf (data) pages first allocated sequentially, sorted by search key; then index pages allocated, and then overflow pages.
- ❖ *Index entries:* <search key value, page id>; they “direct” searches for *data entries*, which are in leaf pages.
- ❖ *Search:* Start at root; use key comparisons to go to leaf.  
 $I/O \text{ cost} \propto \log_F N$  ;  $F = \# \text{ entries/index pg}$ ,  $N = \# \text{ leaf pgs}$
- ❖ *Insert:* Find leaf data entry belongs to, and put it there.
- ❖ *Delete:* Find and remove from leaf; if empty overflow page, de-allocate.




**Static tree structure:** *inserts/deletes affect only leaf pages.*

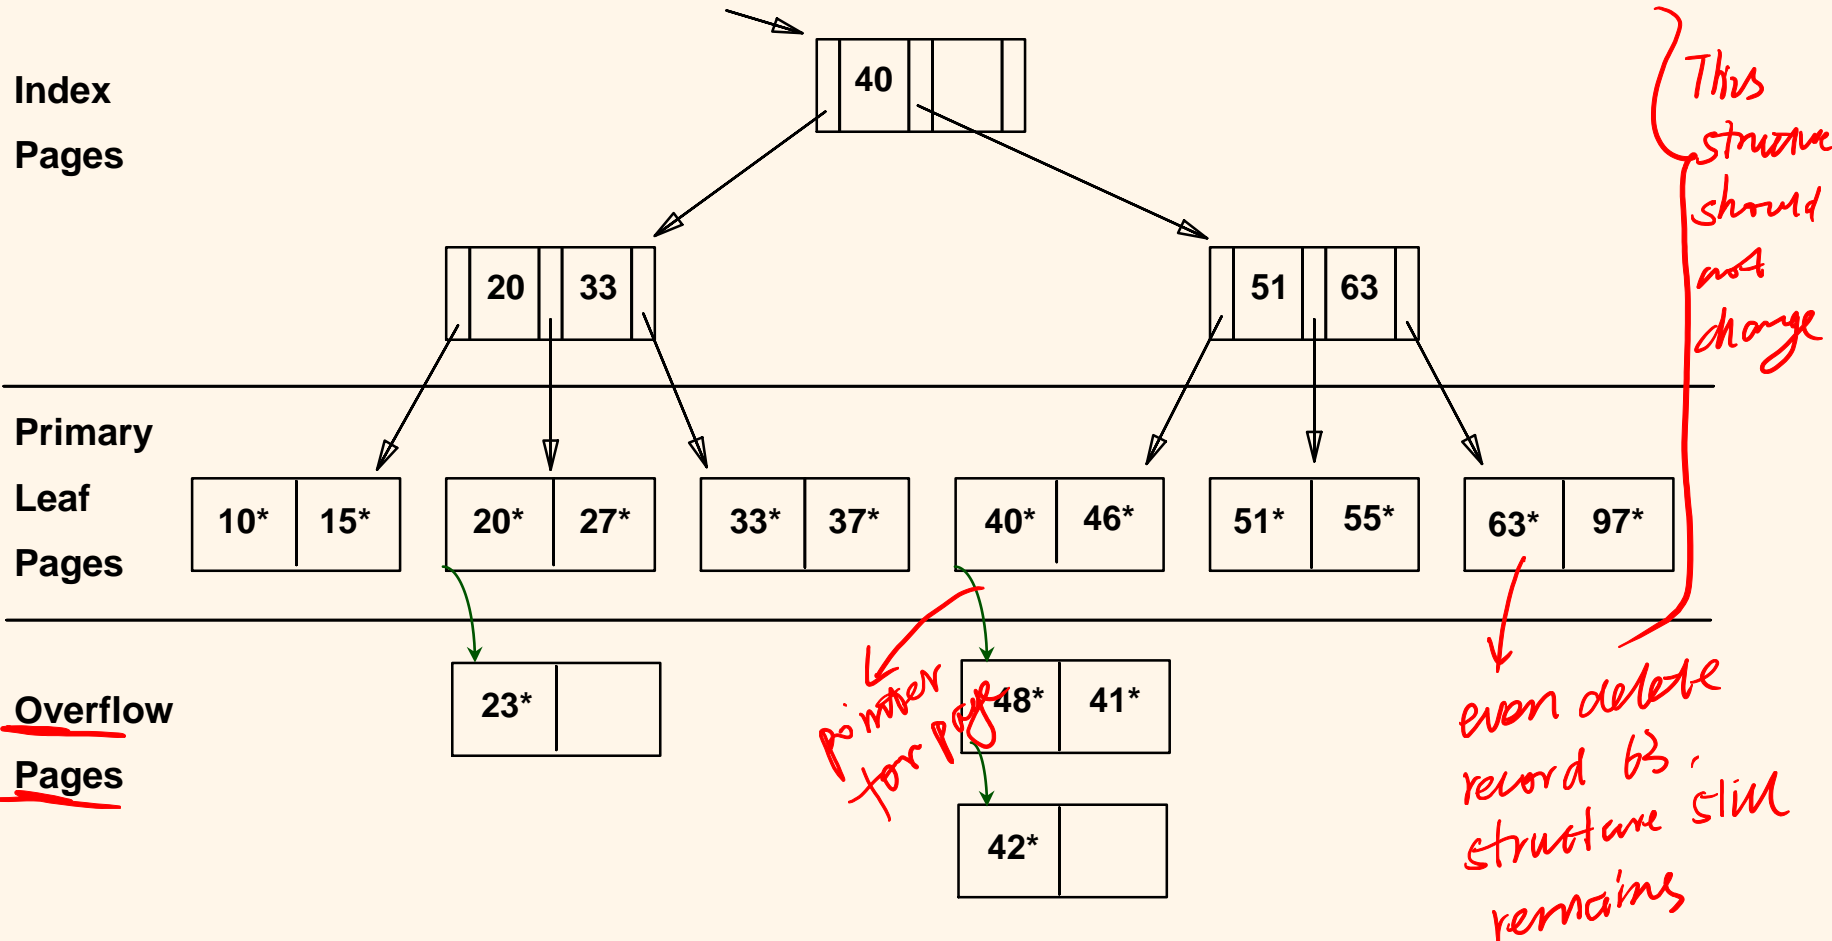


# Example ISAM Tree

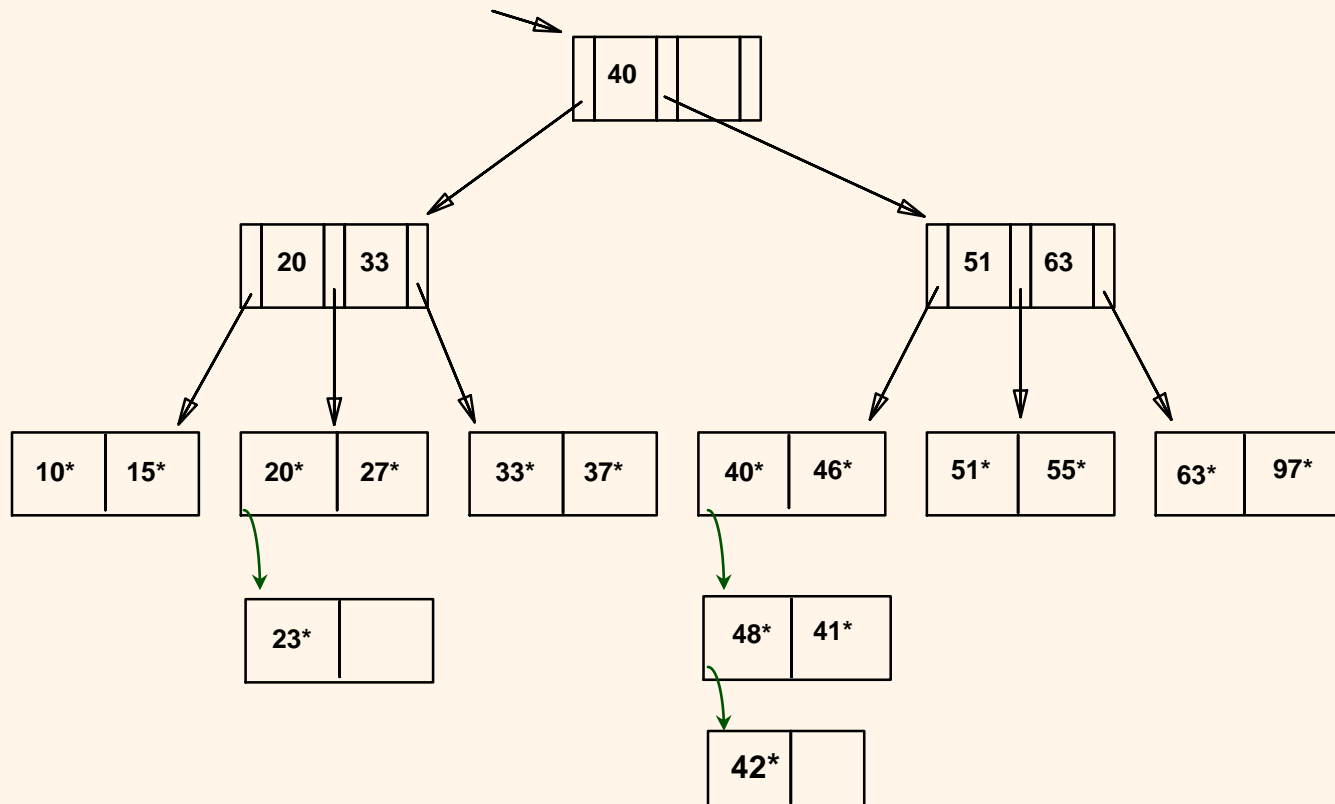
- ❖ Suppose each node can hold 2 entries (really more like 200 since the nodes are disk pages)



# Let's Insert 23\*, 48\*, 41\*, 42\* ...



*... Then Delete 42\*, 51\*, 97\**



*Note how 51\* still appears in index levels, but **not** in leaf!*