

CS222/CS122C: Principles of Data Management

UCI, Fall 2019
Notes #09

External Sorting

Disk based

Instructor: Chen Li


Why Sort?

- ❖ A classic problem in computer science! (😊)
- ❖ Data sometimes requested in sorted order.
 - E.g., find students in decreasing *gpa* order.
- ❖ Sorting is first step in bulk loading B+ tree index.
- ❖ Sorting useful for eliminating *duplicate copies* (i.e., SELECT DISTINCT) in a collection of records. (Why?)
- ❖ Sort-merge join algorithm involves sorting.
- ❖ Problem: sort 1 TB of data with 1 GB of RAM.
 - Q: Why not just use virtual memory?

* Access heap file using sorted RIDs from B⁺ tree.

2-Way Sort: Requires 3 Buffers

- ❖ Pass 1: Read a page, sort it, write it out
 - only one buffer page is used
- ❖ Pass 2, 3, ..., etc.: Read and merge pairs of runs
 - three buffer pages are used:



Sorting $N=2^k$ Pages of Data

❖ Pass 0:

- Read, sort, write $\rightarrow 2^k$ 1-page runs (subfiles)

❖ Pass 1:

- Read+merge 1-page pairs, write $\rightarrow 2^{k-1}$ 2-page runs

❖ Pass 2:

- Read+merge 2-page pairs, write $\rightarrow 2^{k-2}$ 4-page runs

...

❖ Pass $k-1$:

- Read+merge 2^{k-2} -page pairs, write $\rightarrow 2$ 2^{k-1} -page runs

❖ Pass k :

- Read+merge 2^{k-1} -page pairs, write $\rightarrow 1$ 2^k -page result

Two-Way External Merge Sort

❖ Each pass we read + write each page in file.

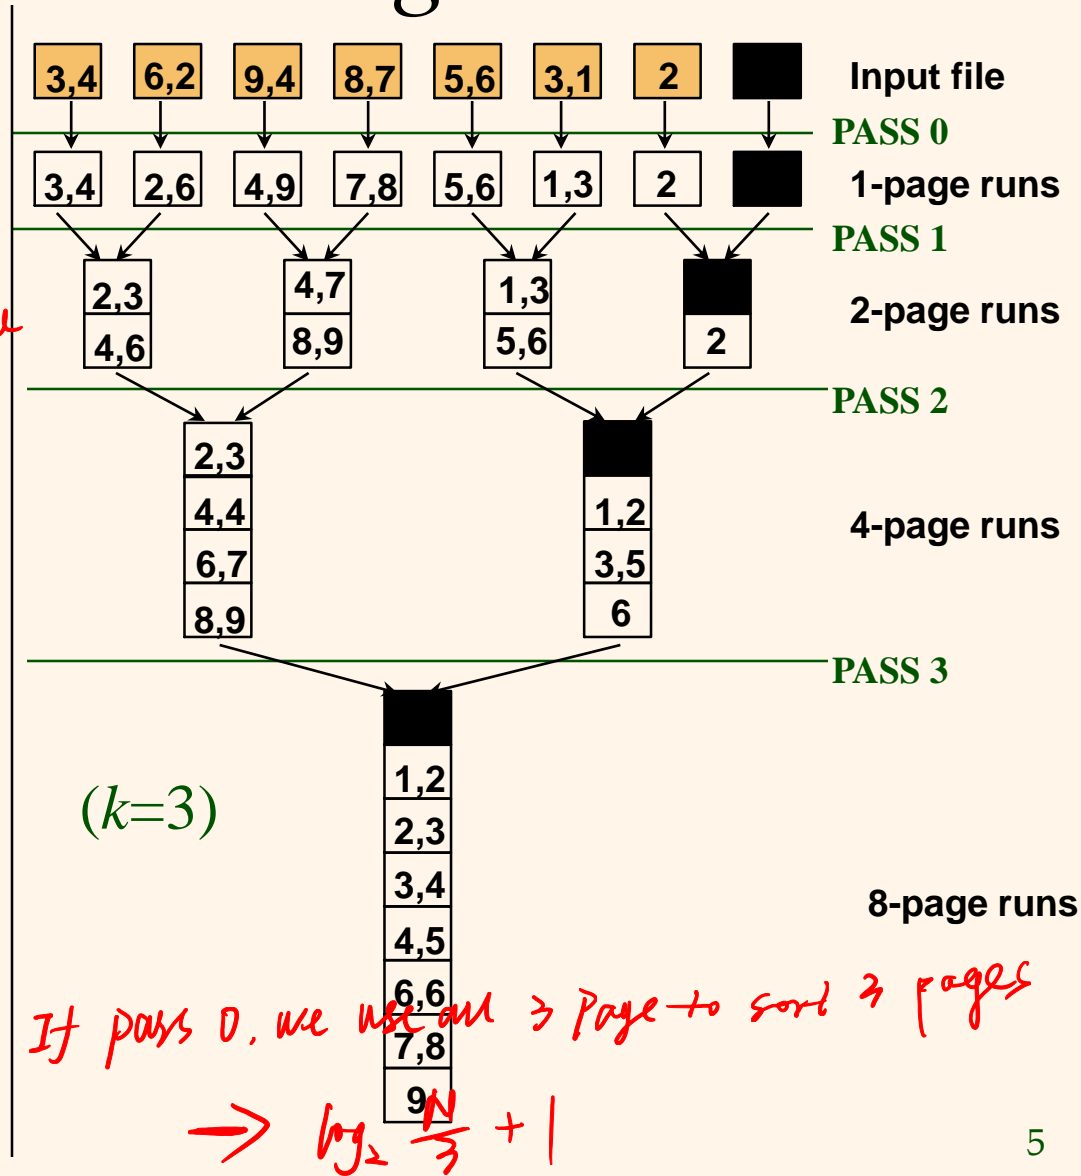
❖ N pages in the file \Rightarrow the number of passes $= \lceil \log_2 N \rceil + 1$ *Sort the page*

❖ So total I/O cost is:

$$2N(\lceil \log_2 N \rceil + 1)$$

❖ Idea: **Divide and conquer**: sort subfiles and merge

❖ Q: See any room to do better, w/just 3 pages?




General External Merge Sort

More than 3 buffer pages. How can we utilize them?

*Read B pages
use quick sort,
and flush back
to disk
No output page is
lost*

❖ To sort a file with N pages using B buffer pages:

- **Pass 0: use B buffer pages.** Produce $\lceil N / B \rceil$ sorted runs of B pages each. (Actually $B-1$ instead of B w/variable-length records.)
- **Pass 2, ..., etc.: merge $B-1$ runs.**



Cost of External Merge Sort

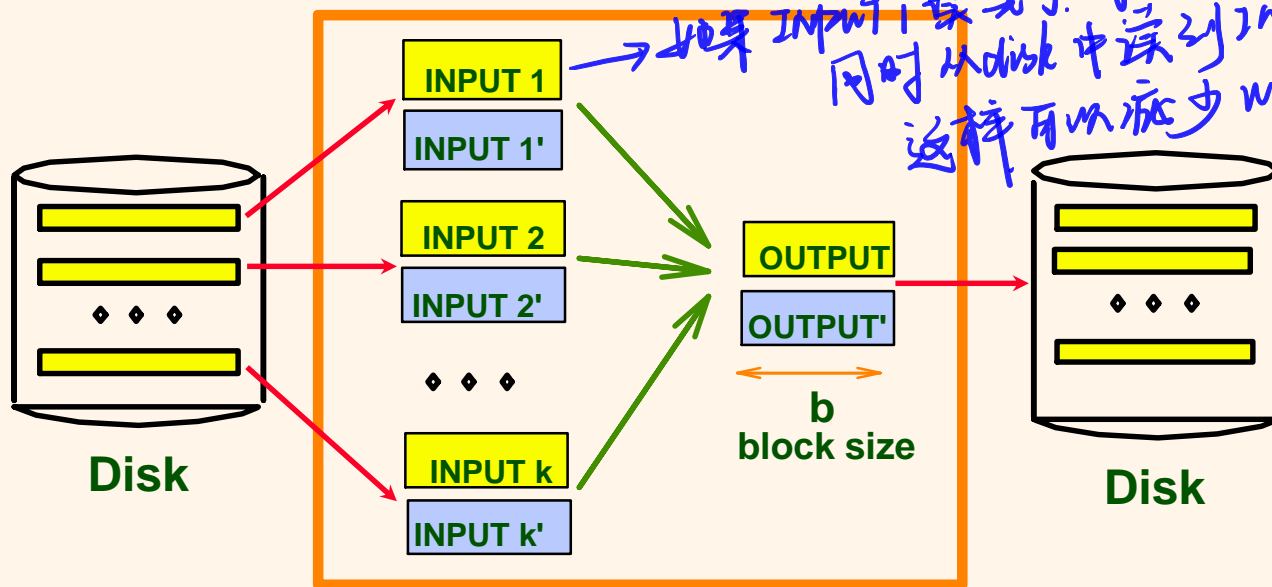
- ❖ Number of passes: $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- ❖ Cost = $2N * (\# \text{ of passes})$
- ❖ E.g., with 5 buffer pages, to sort 108 page file:
 - In memory sort B pages* Pass 0: $\lceil 108 / 5 \rceil = 22$ sorted runs of 5 pages each (last run is only 3 pages)
 - Pass 1: $\lceil 22 / 4 \rceil = 6$ sorted runs of 20 pages each (last run is only 8 pages)
 - Pass 2: 2 sorted runs, 80 pages and 28 pages
 - Pass 3: Sorted file of 108 pages

Number of Passes of External Sort

N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

Double Buffering

- ❖ To reduce wait time for I/O request to complete, can *prefetch* into 'shadow block'.
 - Potentially, more passes; in practice, most files *still* sorted in *2-3* passes.



B main memory buffers, k-way merge

Sorting Summary

- ❖ External sorting is important; DBMS may dedicate part of buffer pool for sorting!
- ❖ External merge sort minimizes disk I/O cost:
 - Pass 0: Produce sorted *runs* of size B (# buffer pages), or of size $B-1$ if we are handling variable-length records.
 - Passes > 0 : *Merge* runs (until just one run is produced).
 - # of runs merged at a time depends on B and *block size*.
 - Larger block size means less I/O cost per page of data.
 - Larger block size means fewer runs merged per step.
 - In practice, # of passes needed rarely more than 2 or 3.

SAL:

called SPJ =

↙
selection
project
join.

select


from

where.

group by

having

order by




$$BR \log(B \cdot R) \cdot \frac{N}{B}$$

in-memory cost for B pages

$$(2 \log(B-1))NR$$

push/pop time for heap



- Build the initial heap
- keep a cursor with one input page
- use pop / push to go through each page.
- whenever the output page is full, flush it to disk
- if one input page is exhausted, read next page

$$\downarrow$$

$$\approx \log(B-1) NR$$

Q: for N page, what's the minimum # of pages that can
sort N pages using 2 iteration

↓
(pass 0, pass 1)
↙ ↘
in-memory sort merge sort

$$\left\{ \begin{array}{l} \frac{N}{B} \leq B-1 \\ \text{or} \\ N \leq B(B-1) \end{array} \right.$$