

CS222/CS122C: Principles of Data Management

UCI, Fall 2019
Notes #12

Set operations, Aggregation, Cost Estimation

Instructor: Chen Li

Relational Set Operations

- ❖ Intersection and cross-product special cases of join.
- ❖ Union (**distinct**) and Except similar; we'll do union.
- ❖ Sorting-based approach to union:
 - Sort both relations (on combination of all attributes).
 - Scan sorted relations and merge them.
 - *Alternative*: Merge runs from Pass 0+ for *both* relations (!).
- ❖ Hash-based approach to union (from Grace):
 - Partition both R and S using hash function $h1$.
 - For each S-partition, build in-memory hash table (using $h2$), then scan corresponding R-partition, adding truly new S tuples to hash table while discarding duplicates.

Sets versus Bags

- ❖ UNION, INTERSECT, and DIFFERENCE (EXCEPT or MINUS) use the set semantics.
- ❖ UNION ALL just combine two relations without considering any correlation.

Sets versus Bags

Example:

- ❖ $R = \{1, 1, 1, 2, 2, 2, 2, 3\}$
- ❖ $S = \{1, 2, 2\}$
- ❖ $R \text{ UNION } S = \{1, 2, 3\}$ ← *eliminate duplicate*
- ❖ $R \text{ UNION ALL } S = \{1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 3\}$
- ❖ $R \text{ INTERSECT } S = \{1, 2\}$
- ❖ $R \text{ EXCEPT } S = \{3\}$

Notice that some DBs such as MySQL don't support some of these operations.

Aggregate Operations (*AVG*, *MIN*, ...)

❖ Without grouping:

- In general, requires scanning the full relation.
- Given an index whose search key includes *all* attributes in the *SELECT* or *WHERE* clauses, can do an index-only scan.

❖ With *grouping*:

- *Sort* on the *group-by* attributes, then scan sorted result and compute the aggregate for each group. (Can improve on this by combining sorting and aggregate computations.)
- Or, similar approach via *hashing* on the group-by attributes.
- Given tree index whose search key includes all attributes in *SELECT*, *WHERE* and *GROUP BY* clauses, can do an index-only scan; if group-by attributes form prefix of search key, can retrieve data entries (tuples) in the group-by order.

Aggregation State Handling

❖ State: `init()`, `next(value)`, `finalize()` → agg. value

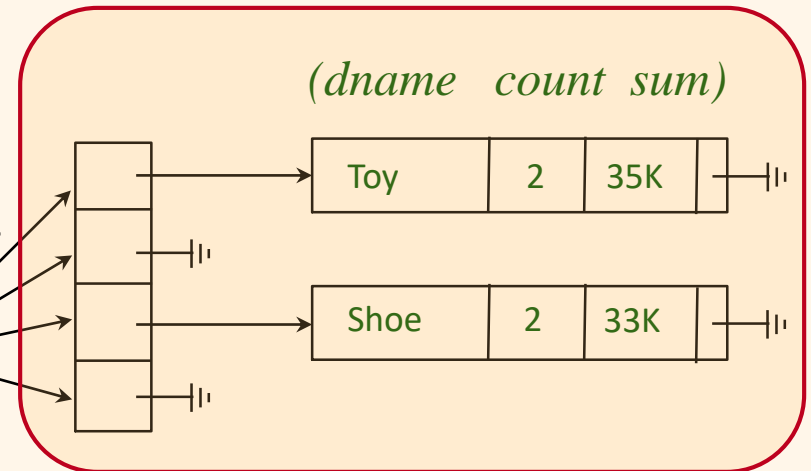
❖ Consider the following query

```
SELECT E.dname, avg(E.sal)
FROM Emp E
GROUP BY E.dname
```

Emp		
ename	sal	dname
Joe	10K	Toy
Sue	20K	Shoe
Mike	13K	Shoe
Chris	25K	Toy
Zoe	50K	Book
...

$h(\text{dname}) \bmod 4$

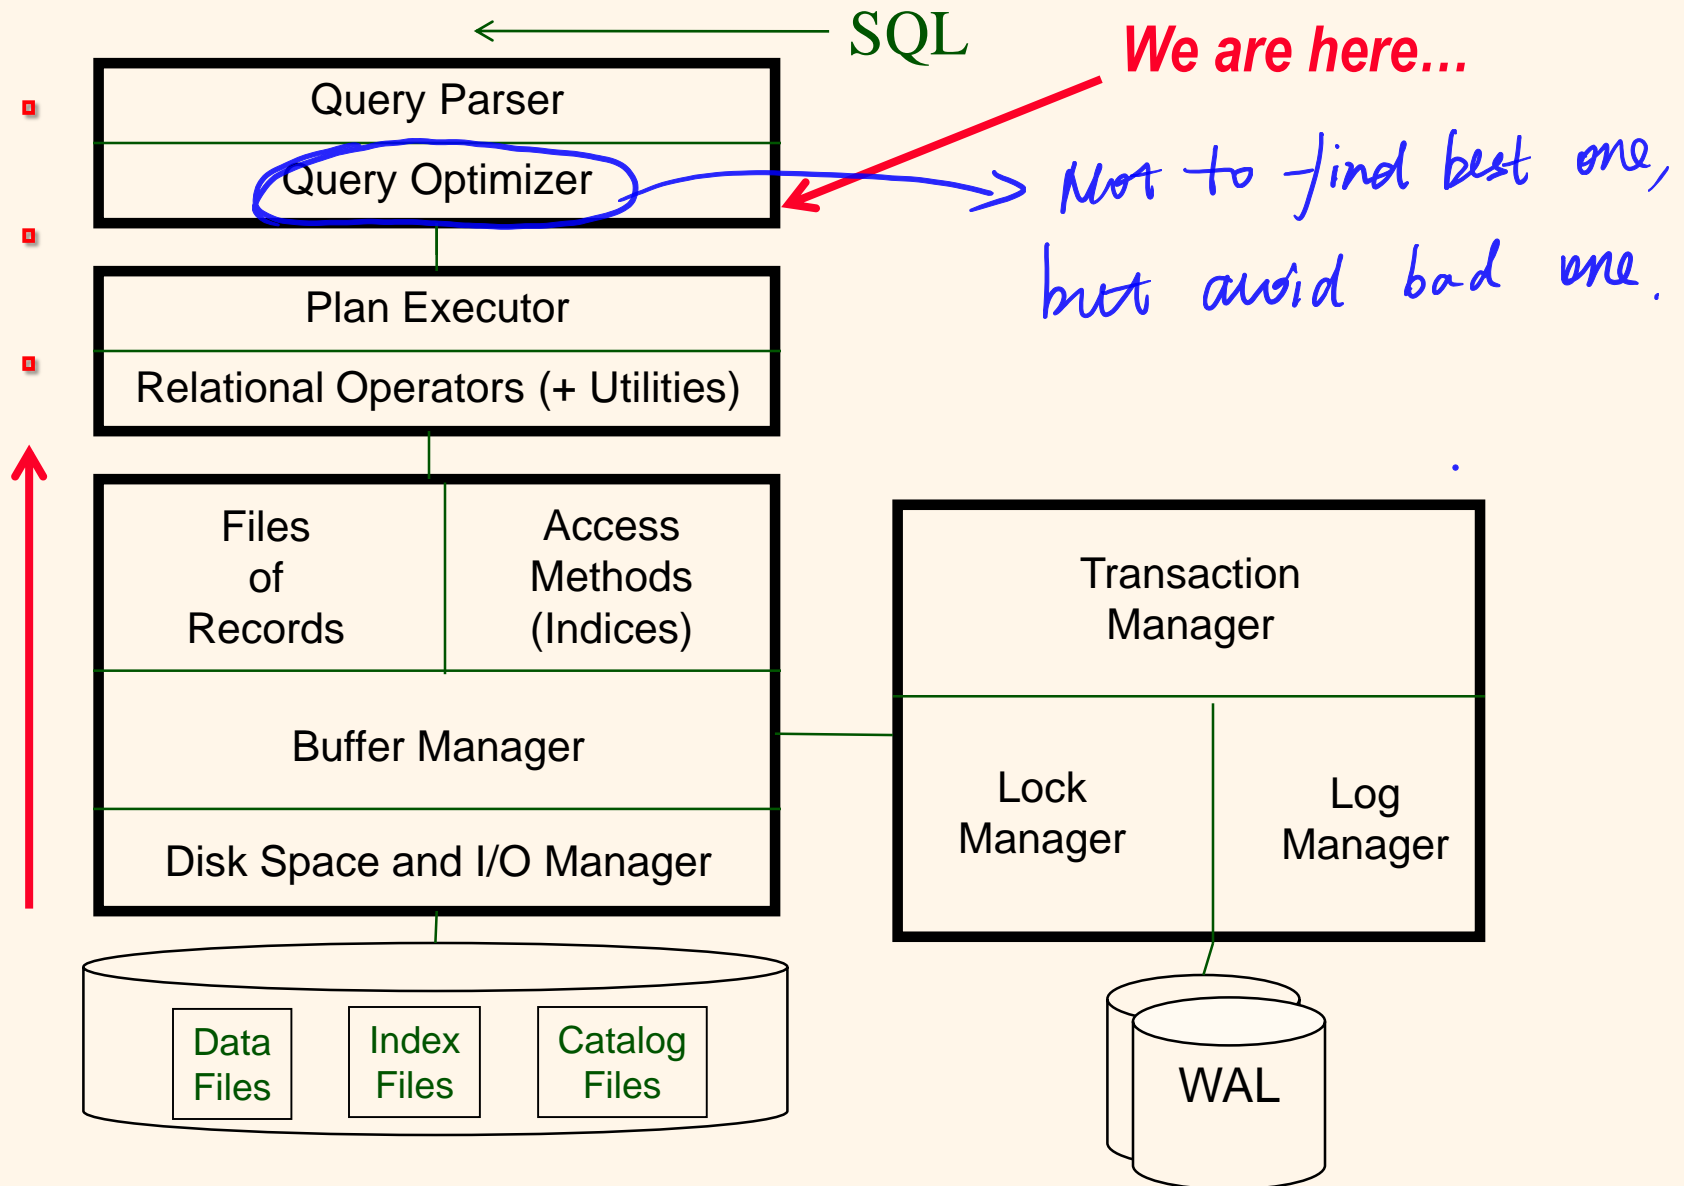
h



Aggregate state (per unfinished group):

- Count: # values
- Min: min value
- Max: max value
- Sum: sum of values
- Avg: count, sum of values

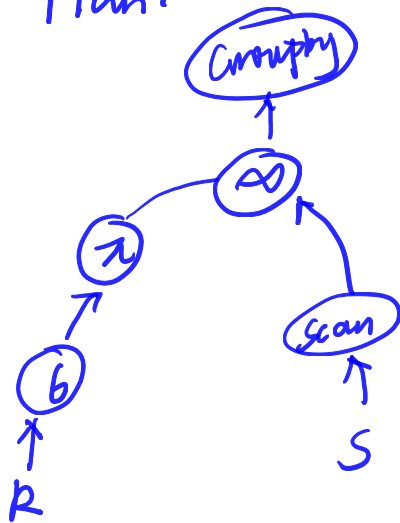
DBMS Structure (from Lecture 1)



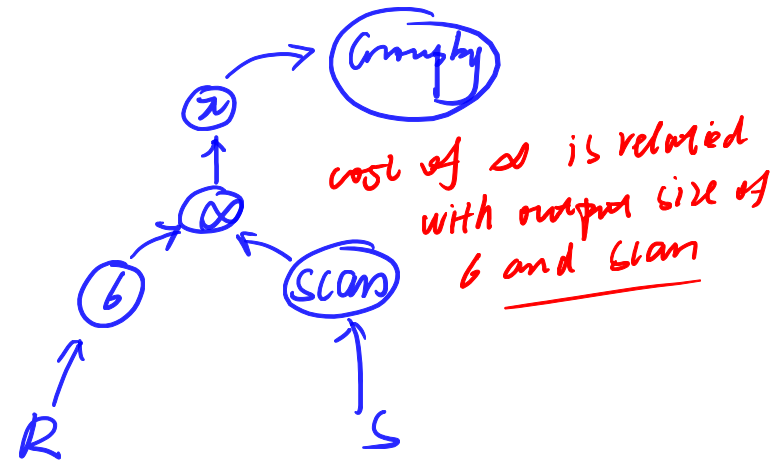
Cost Estimation

- ❖ For each plan considered, need to estimate its cost:
 - Must *estimate cost* of each operation in plan tree.
 - Depends on their input cardinalities.
 - We've already discussed how to estimate the cost of one operation (sequential scan, index scan, join, etc.)
 - Must also *estimate size of result* for each operation in the query plan tree!
 - Use information about the input relations.
 - For selections and joins, assume independence of query predicates.

Plan 1



Plan 2



Costs { Disk IOs
Memory requirements
CPU cycles
Network costs

"Cost function"

Q₁: cost function

Q₂: Estimate cost of query

Q₃: Estimate cost of operator.

the cost of parent node is related to the cardinality of output of child node,

hard to estimate accurately

Estimate the cardinality of an operator.

1) $G =$

→ Range

3) Boolean

4) Sort → doesn't change cardinality

Estimating cost of a selection

- ❖ In our earlier lectures, we already discussed how to estimate the cost of a selection predicate (scan, or using B+ tree)
- ❖ We also analyzed the complexity of those operators (e.g., sort, different join methods)
- ❖ Next we will mainly focus on size estimation of intermediate results

Intermediate Result Size Estimation

- ❖ Optimizers use **statistics** in catalog to estimate the cardinalities of operators' inputs and outputs
 - Simplifying assumptions: uniform distribution of attribute values, independence of attributes, etc.
- ❖ For each relation R:
 - Cardinality of R ($|R|$), avg R-tuple width, and # of pages in R ($||R||$) – pick any 2 to know all 3 (given the page size)
- ❖ For each (indexed) attribute of R:
 - Number of distinct values $|\pi_A(R)|$
 - Range of values (i.e., low and high values)
 - Number of index leaf pages
 - Number of index levels (if B+ tree)

Simple Selection Queries (σ_p)

❖ Equality predicate (p is “ $A = \text{val}$ ”)

- $|Q| \approx |R| / |\pi_A(R)|$ *number of unique value \rightarrow should be in catalog file.*
 - Translation: R's cardinality divided by the number of distinct A values
 - Assumes all values equally likely in R.A (uniform distribution)
 - Ex: SELECT * FROM Emp WHERE age = 23;
- RF (a.k.a. selectivity) is therefore $1 / |\pi_A(R)|$

❖ Range predicate (p is “ $\text{val}_1 \leq A \leq \text{val}_2$ ”)

- $|Q| \approx |R| * ((\text{val}_2 - \text{val}_1) / ((\text{high}(R.A) - \text{low}(R.A)))$
 - Translation: Selected range size divided by full range size
 - Again assumes uniform value distribution for R.A
 - (Simply replace val_i with high/low bound for a one-sided range predicate)
 - Ex: SELECT* FROM Emp WHERE age \leq 21;
- RF is $((\text{val}_2 - \text{val}_1) / ((\text{high}(R.A) - \text{low}(R.A)))$

Boolean Selection Predicates

❖ Conjunctive predicate (p is “ p_1 and p_2 ”)

- $RF_p \approx RF_{p_1} * RF_{p_2}$

- Ex 1: SELECT * FROM Emp WHERE age = 21 AND gender = 'm'
- Assumes independence of the two predicates p_1 and p_2 (uncorrelated)
- Ex 2: SELECT * FROM Student WHERE major = 'EE' AND gender = 'm'

❖ Negative predicate (p is “not p_1 ”)

- $RF_p \approx 1 - RF_{p_1}$

- *Translation:* All tuples minus the fraction that satisfy p_1
- Ex: SELECT * FROM Emp WHERE age \neq 21

❖ Disjunctive predicate (p is “ p_1 or p_2 ”)

- $RF_p \approx RF_{p_1} + RF_{p_2} - (RF_{p_1} * RF_{p_2}) \leftarrow Q: \text{Why?}$

- *Translation:* All (unique) tuples satisfying either predicate
- Ex: SELECT * FROM Student WHERE major = 'EE' OR gender = 'm'

Two-way Equijoin Predicates

Join.


- ❖ Query Q: $R \text{ join } S \text{ on } R.A = S.B$
- ❖ Assume join value set containment (FK/PK case)
 - $\pi_A(R)$ is a subset of $\pi_B(S)$ or vice versa
 - Translation: “Smaller” relation’s join value set is subset of “larger” relation’s join value set (where “size” is based on # unique values)
 - Ex: SELECT * FROM Student S, Dept D WHERE S.major = D.deptname
 - $|Q| \approx (|R| * |S|) / \max(|\pi_A(R)|, |\pi_B(S)|)$
 - Ex: 100 D.deptname values but 85 S.major values, 10,000 students
 - Estimated size of result is **1/100** of the cartesian product of Student and Dept
 - Again making a uniformity assumption (i.e., about students’ majors)
 - I.e., RF is $1 / \max(|\pi_A(R)|, |\pi_B(S)|)$
- ❖ Repeat same principle to deal with N-way joins

← Q: Why? (Why not 1/85?)

Improved Estimation: Histograms

- ❖ We have been making simplistic assumptions
 - *Specifically:* uniform distribution of values
 - This is definitely violated (all the time ☺) in reality
 - Violations can lead to huge estimation errors
 - Huge estimation errors can lead to Very Bad Choices
- ❖ By the way:
 - In the absence of info, System R assumed $1/3$ and $1/10$ for range queries and exact match queries, respectively
 - (Q: What might lead to the absence of info?)
- ❖ How can we do better in the OTHER direction:
 - Keep track of the most and/or least frequent values
 - Use histograms to better approximate value distributions

Equi-width Histograms



- ❖ Divide the domain into B buckets of equal width
 - E.g., partition Kid.age values into buckets
- ❖ Store the bucket boundaries and the sum of frequencies of the values with each bucket

Histogram Construction

❖ Initial Construction:

- Make one full one pass over R to construct an accurate equi-width histogram
 - Keep a running count for each bucket
- If scanning is not acceptable, use sampling
 - Construct a histogram on R_{sample} , and scale the frequencies by $|R|/|R_{sample}|$


❖ Maintenance Options:

- Incremental: for each update to R , increment or decrement the corresponding bucket frequencies (Q : Cost?)
- Periodically recompute: distribution changes slowly!

Using Equi-width Histograms

❖ Q: $\sigma_{A=5}(R)$


- 5 is in bucket $[5,8]$ (with 19 tuples)
- Assume uniform distribution within the bucket
- Thus $|Q| \approx 19/4 \approx 5$.
- Actual value is 1



❖ Q: $\sigma_{A \geq 7 \ \& \ A \leq 16}(R)$

- $[7,16]$ covers $[9,12]$ (27 tuples) and $[13,16]$ (13 tuples)
- $[7,16]$ partially covers $[5,8]$ (19 tuples)
- Thus $|Q| \approx 19/2 + 27 + 13 \approx 50$
- Actual value = 52.

Equi-height Histogram *adaptive to skew*



- ❖ Divide the domain into B buckets with (roughly) the same number of tuples in each bucket
- ❖ Store this number and the bucket boundaries
- ❖ Intuition: high frequencies are more important than low frequencies


Histogram Construction

❖ Construction:

- Sort all R.A values, and then take equally spaced slices
 - *Ex:* 1 2 2 3 3 5 6 6 6 6 6 6 7 7 7 7 8 8 8 8 8 ...
- Sampling also applicable here

❖ Maintenance:

- Incremental maint
 - Split/merge buckets (B+ trees)
- Periodic recomputation




Using an equi-height histogram

❖ Q: $\sigma_{A=5}(R)$

- 5 is in bucket [1,7] (with 16 tuples)
- Assume uniform distribution within the bucket
- Thus $|Q| \approx 16/7 \approx 2.3$. (actual value = 1)

❖ Q: $\sigma_{A \geq 7 \ \& \ A \leq 16}(R)$

- [7,16] covers [8,9],
- [7,16] partially covers [10,11],
- Thus $|Q| \approx 16/7 + 16 + 16 + 16 \approx 50$
- Actually $|Q| = 52$.



Can Combine Approaches

- ❖ If values are badly skewed
 - Keep high/low frequency value information in addition to histograms
 - Could even apply the idea recursively: keep this sort of information for each bucket
 - “Divide” by converting values to some # ranges
 - “Conquer” by keeping some statistics for each range
- ❖ Some “statistical glitches” to be aware of
 - Parameterized queries
 - Runtime situations

Histogram \rightarrow distribution